

The CHRO Lens

An audience-specific reading of the OIA Reference Model for chief human resources officers and the people organisations they lead.

By Mark Templeton · Part of the OIA Lens System · Draft v0.1 · 2026 · Published by the Agentics Foundation

Anticipated lens — draft for author review. The Lens System v0.3 names the CHRO Lens as a future lens to be published “as practitioner engagement and demand support,” and observes that workforce transformation is the half of the AI conversation least well served by existing frameworks, because those frameworks were built for human workforces only. This draft is authored to the Lens System methodology (5–7 interpretive lenses across 2–3 orthogonal axes) ahead of formal publication; it should be reviewed before it is treated as canonical.

1. What this document is

The OIA Reference Model describes the structural shape of enterprise intelligent systems — ten layers, seven cross-layer spans, three defining properties, four named threat surfaces — vendor-neutral and audience-neutral. People leaders do not work from structural specifications; they work from workforce strategy, role design, and governance built for humans. Yet agents now act inside the enterprise as a non-human workforce — and the existing HR frameworks do not see them. Without translation, the OIA never reaches the workforce conversation it should inform.

This is the translation. The CHRO Lens reads the same OIA architecture through the questions a people leader is actually asking about a workforce that is now part human, part agent. The architecture is invariant; the lens changes the language. Layer 7, Layer 9, and the Identity span together describe the

architectural

surface at which workforce policy for agents is implemented.

2. What is in the CHRO planning conversation in 2026

- **Augmentation versus displacement** is the defining workforce question, and it is being decided team by team.
- **Agents are a non-human workforce** operating alongside employees, largely ungoverned by people policy.
- **The employee lifecycle does not yet extend to agents** — there is no standard onboarding, scoping, baselining, or offboarding for them.
- **Performance management in mixed teams** (humans + agents) has no established practice.
- **Reskilling at scale** must keep pace with capability that changes monthly.
- **Segregation of duties and least privilege** — human governance concepts — now apply to agents too.
- **The “three actors” reality** (humans, agents, foundation models) reframes “workforce” itself.

3. The six interpretive lenses on three axes

Workforce axis · What AI is to people

AI as Augmentation. *Design roles where agents extend people, not replace judgement.* Anchor: job-architecture / capability frameworks.

AI as Non-Human Workforce. *Govern agents with the discipline applied to employees.*

Onboarding, scoping,

behavioural baselining, offboarding — the agent lifecycle. Anchor: workforce-governance practice.

Governance axis · How the mixed workforce is run

AI as Managed Headcount. *Bring agents into role design, access, and segregation of duties.*

Every agent

is, in effect, an employee with elevated privileges. Anchor: HR governance + NIST CSF (Govern).

AI as Performance-Managed. *Measure and manage mixed human-agent teams.* Anchor:

performance-management

frameworks.

Capability axis · How the workforce adapts

AI as Reskilling Imperative. *Keep human capability ahead of automation, continuously.* Anchor: learning & development frameworks.

AI as Culture and Trust. *Sustain trust and accountability in a part-agent workforce.* Anchor: change-management / organisational-trust practice.

4. The four threat surfaces as workforce exposures

- **Mimicry** (Layer 9) — agents impersonating people (and vice versa) undermines trust and accountability in mixed teams.
- **Agent-in-the-middle** (Layer 7) — an ungoverned agent-tool boundary is the workforce equivalent of an unsupervised employee with unchecked access.
- **Agent spoofing** (Identity span) — an agent inheriting authority it was never granted is a segregation-of-duties failure.
- **Model compromise** (Layer 5) — a workforce of agents on a silently changed model is a workforce whose behaviour drifted without a performance review.

5. The layers and spans through the CHRO Lens

Workforce policy for agents concentrates at a few points in the architecture.

- **Layer 9 · Human and Browser Interface** — where humans and agents meet; the surface for trust, disclosure, and human-agent teaming.
- **Layer 7 · Orchestration and Workflow** — where agents *act*; the locus of agent role, scope, and supervision.
- **Identity span** — the three actors (humans, agents, foundation models); the mechanism for onboarding, scoping, baselining, and offboarding an agent workforce.
- **Auditability span** — the evidence of who (or what) did what, essential to performance management and

accountability in mixed teams.

- Elsewhere (Layers 0–6, Layer 8) the CHRO’s interest is indirect, mediated through the technology and security organisations.

6. Adopting the lens

- **As a workforce-governance lens** — extend the employee lifecycle to agents using the Identity span and Layers 7 and 9 as the architectural anchor.
- **As a role-design vocabulary** — reason about augmentation versus displacement against a structural map.
- **As a cross-functional bridge** — align HR, security (“every agent is an employee with elevated privileges”), and technology on the same architecture.

The Foundation invites the people-leadership community to engage with the lens, push back on what does not fit, and contribute to its refinement.