

The General Counsel and Compliance Lens

An audience-specific reading of the OIA Reference Model for general counsel, chief compliance officers, and the legal and regulatory functions they lead.

By Mark Templeton · Part of the OIA Lens System · Draft v0.1 · 2026 · Published by the Agentics Foundation

Anticipated lens — draft for author review. The Lens System v0.3 names the General Counsel and Compliance Lens as a future lens to be published “as practitioner engagement and demand support.” This draft is authored to the Lens System methodology (5–7 interpretive lenses across 2–3 orthogonal axes) ahead of formal publication; it should be reviewed before it is treated as canonical.

1. What this document is

The OIA Reference Model describes the structural shape of enterprise intelligent systems — ten layers, seven

cross-layer spans, three defining properties, four named threat surfaces — vendor-neutral and audience-neutral. Legal and compliance leaders do not work from structural specifications; they work from

liability, regulatory exposure, and contractual surface, with the regimes evolving in real time.

Without

translation, the OIA never reaches the accountability conversation it should inform.

This is the translation. The General Counsel and Compliance Lens reads the same OIA architecture through the

questions a legal leader is actually asking — where liability attaches, where regulatory duties bind, and

where the evidence to defend them physically lives. The architecture is invariant; the lens changes the

language. The OIA’s cross-layer spans — auditability, provenance, sovereignty, identity — are precisely the

artefacts that legal accountability and regulatory compliance reduce to.

2. What is in the legal/compliance planning conversation in 2026

- **Accountability for agent action is unsettled** — system owner, model provider, vendor, or the enterprise?
- **The regulatory map is moving under the deployment.** EU AI Act, US state-level AI law, and sectoral rules apply with extraterritorial reach and shifting classifications.
- **Training-data IP and output responsibility are open exposures** — provenance of what went in, ownership of what comes out.
- **Vendor contracts were written for software, not for actors.** Indemnities, warranties, and liability caps do not map to autonomous behaviour.
- **Data residency and sovereignty are compliance obligations**, not just architecture.
- **Disclosure and record-keeping duties assume reconstructability** the systems may not provide.
- **Regulators are asking for human oversight** the enterprise must be able to locate and evidence.

3. The six interpretive lenses on three axes

Liability axis · Who is answerable

AI as Liable Actor. *Locate where responsibility for an agent's action attaches.* Autonomy at Layer 7

concentrates the question of accountability. Anchor: tort / agency doctrine.

AI as Output Owner. *Establish responsibility for model output and its IP.* Provenance from Layer 3 to

Layer 9 governs both. Anchor: IP and product-liability frameworks.

Regulatory axis · What binds the enterprise

AI as Regulated System. *Classify high-risk systems and locate human oversight.* Anchor: **EU AI Act**

(and US state-level AI law).

AI as Managed Compliance. *Attest against an AI management system.* Anchor: **ISO/IEC 42001**,

NIST AI RMF.

Evidence axis · What can be proven

AI as Provable Record. *Ensure the record needed to defend a claim exists by design.* Auditability cannot

be retrofitted; it originates at Layer 3. Anchor: e-discovery / records-management duties.

AI as Sovereign Obligation. *Meet data-residency and jurisdictional duties as compliance, not preference.*

Established at Layer 2, threading upward. Anchor: data-protection and sectoral regimes.

4. The four threat surfaces as legal exposures

- **Mimicry** (Layer 9) — impersonation and synthetic-content fraud; disclosure and consumer-protection exposure.
- **Agent-in-the-middle** (Layer 7) — an un-auditable intermediary undermines the evidentiary record.
- **Agent spoofing** (Identity span) — inherited authority creates unbounded and unattributable liability.
- **Model compromise** (Layer 5) — reliance on a silently mutable model is a warranty, disclosure, and contractual-risk exposure.

5. The ten layers through the General Counsel Lens

For each layer: where liability, regulatory duty, and evidentiary obligation physically live.

- **Layer 9 · Human and Browser Interface** — consumer protection, disclosure, impersonation liability.
- **Layer 8 · Continuity Fabric** — records integrity, spoliation risk, post-incident reconstructability.
- **Layer 7 · Orchestration and Workflow** — agent-action liability; the locus of the accountability question.
- **Layer 6 · Context and Knowledge** — IP in retrieved corpora; confidentiality and privilege exposure.
- **Layer 5 · Inference and Retrieval** — model-output responsibility; vendor warranty and CVE disclosure.
- **Layer 4 · Model Training and Adaptation** — training-data IP, licensing, and provenance

obligations.

- **Layer 3 · Agent Data Substrate** — privacy, data-subject rights, and the origination of auditability.
- **Layer 2 · Sovereign Infrastructure** — data residency, jurisdiction, and cross-border transfer duties.
- **Layer 1 · Silicon Abstraction** — attested-execution as evidence of controls (Verifiability span).
- **Layer 0 · Physical Compute** — supply-chain and export-control exposure.

6. Adopting the lens

- **As a compliance-mapping lens** — locate, layer by layer and span by span, where regulatory and contractual exposure physically lives.
- **As a contracting vocabulary** — negotiate indemnities, warranties, and liability caps against the architecture rather than the feature list.
- **As a disclosure and evidence anchor** — ensure the auditability and provenance needed to defend a claim are designed in, not promised later.

The Foundation invites the legal and compliance community to engage with the lens, push back on what does not fit, and contribute to its refinement.