

# The OIA Lens System

*Three lenses that make the OIA Reference Model usable for the executive audiences accountable for AI infrastructure, applications, and oversight.*

**An integral part of the OIA Reference Model · Version 0.3 · 2026 · Published by the Agentics Foundation**

---

## Executive Summary

---

The OIA Reference Model describes the architecture of enterprise intelligent systems — ten layers, seven cross-layer spans, written in plain structural terms. That neutrality is what gives the architecture its durability. It also leaves a gap: most executives accountable for AI do not work from structural specifications. They work from their own frameworks, their regulatory environment, and the questions on their desk this quarter.

The OIA Lens System closes that gap. It is published as an integral part of the OIA Model and produces audience-specific overlays that translate the architecture into the language of the executives who must apply it — without changing the architecture itself. Three lenses are published with this release, calibrated to the executive audiences most directly accountable for AI infrastructure, applications, and oversight.

### LENS 1 · THE CISO LENS

**The question it answers.** What does it take to keep the enterprise secure when humans, agents, and the foundation models themselves are now all actors that need continuous trust governance?

**Why it matters.** AI has changed both the speed of attack and the surface area of who can be attacked. Existing security programs were not architected to govern agents at machine speed. The gap between agent capability and agent governance is where the next breaches will happen.

### LENS 2 · THE CIO/CTO LENS

**The question it answers.** How do we deploy AI into the enterprise without inheriting a generation of architectural debt — and what is the architecture, not the vendor list?

**Why it matters.** Capital is being committed under pressure and vendor decisions are being made before the architectural picture is clear. Without architectural clarity, the result is fragmented, expensive, and locked in to choices made before the implications were understood.

## LENS 3 · THE BOARD AND RISK LENS

**The question it answers.** How does the board govern AI strategy and risk concurrently, with the structural vocabulary required to ask the right questions of management — without becoming technologists?

**Why it matters.** AI risk has moved to the board level faster than any technology category in recent memory. Disclosure regimes are evolving in real time and audit committees lack a structural vocabulary for what to audit. The board needs an architectural anchor for accountability that does not require directors to become technologists.

---

### 1. The Lens System and the OIA Model

---

The OIA Reference Model is intentionally neutral. It describes the architecture of enterprise intelligent systems in plain structural terms — ten layers, numbered 0 through 9, and seven cross-layer spans, written so that a CISO, a board director, and a CFO can each read the same model without it being slanted toward any one of them. That neutrality is what gives the architecture its durability.

It also creates a problem. Reference architectures, in their pure form, are read by very few people. Most executives accountable for AI in their organisations do not work from structural specifications. They work from their own frameworks, their regulatory environment, the questions their boards are asking, and the decisions on their desk this quarter. Without a translation layer, the OIA risks being filed as a security document by the first CISO who picks it up, or as a strategy paper that never reaches the engineers who will build to it. Either outcome leaves the architecture short of its purpose.

The Lens System solves this. It is a set of audience-specific overlays — Lens Maps — published as an integral part of the OIA Model. Each lens reads the same ten layers and seven cross-layer spans, but through the questions, the risks, the framework anchors, and the interpretive frames of a specific audience. Three lenses are being published with this release: the CISO Lens for the security organisation, the CIO/CTO Lens for the technology organisation, and the Board and Risk Lens for directors and risk officers. Together they cover the executive audiences most directly accountable for AI infrastructure, applications, and oversight.

The Lens System is part of the OIA Model, not adjacent to it.

Two points of emphasis. First, the Lens System is part of the OIA Model, not adjacent to it. The architecture and the lenses are designed and published together, on the principle that a reference architecture without a clear path into the audiences who must apply it is incomplete. Second, the

lens approach is open. The Foundation is publishing three lenses with this release; additional lenses for further audiences — chief financial officers, general counsel, chief human resources officers, others — can be added over time, by the Foundation or by the broader practitioner community working under the same approach.

This document explains why the Lens System exists, what it delivers, what each of the three published lenses contains, and where the lens approach can go from here. It is written for the executives, practitioners, and Foundation contributors who will use the Lens System in their own work.

## **2. What the Lens System Delivers**

---

The Lens System delivers three things that an architecture alone cannot. Each is meaningful for a different stakeholder — the executive trying to engage with the OIA, the enterprise trying to align its leadership team, and the architecture itself.

### **2.1 A path into the model for every executive**

Each lens speaks the language of its audience. A CISO who reads the CISO Lens encounters the OIA through the threat surfaces, the framework references, and the planning questions that already structure their work. A CIO who reads the CIO/CTO Lens encounters it through the operational and architectural frameworks they already manage. A board director reading the Board and Risk Lens encounters it through accountability and disclosure frameworks. None of these audiences has to translate the OIA into their own working vocabulary; the lens does that work for them. The barrier to engagement drops from ‘read a reference architecture’ to ‘read a document organised around the questions you are already asking.’

### **2.2 Cross-functional alignment on AI infrastructure**

The harder problem in most enterprises is not getting any one executive to engage with AI architecture — it is getting the whole leadership team aligned on the same architectural picture. AI investment, AI governance, AI risk, and AI workforce decisions all sit at the intersection of multiple executive desks, and most of those conversations today happen without a shared structural reference.

The Lens System changes that. When the CISO, the CIO, and the board are reading the OIA through their respective lenses, they are reading the same architecture. The same ten layers carry weight differently for each of them, but the layers themselves are stable. Decisions about AI infrastructure investment, governance posture, and risk exposure can reference shared structural ground. Cross-functional conversation across executive boundaries becomes architecturally

tractable, not because every executive becomes fluent in every other executive's vocabulary, but because every executive can point at the same layer of the same model and be understood.

## **2.3 Validation for the architecture itself**

Every lens is a stress-test. A reference architecture that bends or breaks under the weight of a specific audience's working vocabulary is not fit for purpose. A reference architecture that holds its shape across the security audience's threat lens, the technology audience's operational lens, and the governance audience's accountability lens is, by that fact, drawn against the structural shape of the problem rather than against any single role's perspective.

The lens work makes one observation about the OIA especially clear: Layer 8 — the Continuity Fabric — emerges as a universal gap across every audience we have engaged with, in their own vocabulary. The security organisation feels its absence as a recovery and rollback gap. The technology organisation feels it as a state-coherence gap, where no single agent framework can be allowed to own continuity for the enterprise. The governance audience feels it as an audit-chain integrity gap, where the model that decided is no longer the model deployed. One layer, three different audience readings, the same architectural absence. That convergence is the strongest possible signal that the OIA is naming the right thing, in the right place, at the right level of generality.

The audiences finding themselves on the page — and finding the same page — is what foundational means.

Together, these three deliverables make the case for the Lens System as more than a companion: as the path by which the OIA earns the foundational position the architecture is designed to hold.

## **3. The Three Published Lenses**

---

The first release of the OIA Lens System publishes three lenses for the executive audiences most directly accountable for AI inside the enterprise. Each lens is calibrated to its audience: the questions they actually ask in their planning conversations, the frameworks already in their working vocabulary, and the parts of the OIA that resonate most strongly with the work in front of them. The treatments below describe what each lens contains and what it does for the audience it serves.

### **LENS 1 · THE SECURITY ORGANISATION**

#### **3.1 The CISO Lens**

The security audience has been carrying AI risk inside the enterprise since well before AI became a

board-level conversation. They have lived the pace change. Agents now act inside the enterprise at machine speed. Deepfakes can clone an executive's voice convincingly enough to authorise a wire transfer. Agent-in-the-middle attacks, where a rogue MCP server inserts itself between a trusted agent and a real tool, are no longer theoretical. The CISO Lens organises the OIA around how this audience already structures its work.

What we hear in CISO planning conversations in 2026 is concrete and urgent. Can we tell humans from agents at every entry point? Are agent identities scoped, attestable, and prevented from inheriting trust they should not have? Can we trust the chain between every agent and every tool, including MCP servers, at machine speed? Are we continuously scoring the trustworthiness of every model we use against the threat reality of this week, not last quarter? How do we recover from an autonomous agent that mis-acted? The mental model converging across the audience is that there are now three actors inside the enterprise that need continuous trust governance — humans, agents, and the foundation models themselves.

The frameworks already in their plans are extensive. NIST CSF organises their cyber programme. NIST AI RMF organises their AI risk governance. MITRE ATLAS provides threat vocabulary for AI-specific adversary techniques, and MITRE D3FEND provides the defender side. The OWASP Top 10 for LLM Applications enumerates the risks they are mapping into existing controls. The Confidential Computing Consortium provides the language for hardware-attested execution as a trust model alternative to vendor promises. None of these is a structural architecture — they describe activities, threats, controls, and management systems. The OIA fills the structural gap they sit on top of.

The OIA's ten layers and seven spans map cleanly to how security teams already organise their thinking. Layer 9 for the human and browser interface, where mimicry and deepfakes land. Layer 7 for the orchestration and tool-call surface, where agent-in-the-middle attacks live. Layer 5 for inference and model trust. Layer 3 for the data substrate, where DLP and shadow-AI exfiltration meet. The cross-layer spans — security, identity, auditability, sovereignty, provenance — are the connective tissue that has been missing as a named architectural concern.

Two parts of the OIA resonate especially with what we hear from CISOs. First, Layer 8 — the Continuity Fabric. CISOs feel its absence as a recovery and rollback gap they cannot articulate in any existing framework: there is no playbook for what to do when an autonomous agent has acted, the action was wrong, and the audit chain is broken because the model that decided is no longer the model deployed. The OIA names this gap. Second, the OIA names agent identity as a category-of-one open question — evidence that the model is engaging with what the security audience encounters in the field, not retrofitted from older frameworks.

The CISO Lens organises this material in a way the audience can pick up and recognise immediately. A question per layer in CISO vocabulary. The four named threat surfaces of the

agentic era — mimicry, agent-in-the-middle, agent spoofing, model compromise — overlaid as failure paths against the layers where they originate and propagate. Seven interpretive lenses on three axes: a threat axis (AI as Amplifier, Domain, Attacker), an operational axis (AI as Defender, Asset, Workforce), and a foundational axis (AI as Trustless). Every layer is anchored to the existing frameworks the audience already maintains in their program plans. The result is a document the security organisation can read in one sitting and use in the next planning meeting.

## **LENS 2 · THE TECHNOLOGY ORGANISATION**

### **3.2 The CIO/CTO Lens**

The technology executive has a different problem than the CISO. They are not primarily worried about being attacked — they are worried about deploying AI into the enterprise without inheriting a generation of architectural debt. Capital is being committed under pressure. Vendor decisions are being made before the architectural picture is clear. The board is asking for an AI strategy and the CIO is the one expected to articulate it in terms the board can govern. The CIO/CTO Lens is calibrated to that conversation.

The questions in CIO and CTO planning conversations in 2026 are about coherence, not capability. The capability is largely available; the question is how to assemble it. Where does our stack end and our partners' stack begin — and is that boundary the same boundary our procurement team thinks it is? What is our exposure to GPU concentration, and is that exposure architectural or just contractual? Where do we have de facto multi-cloud AI deployment without a de jure multi-cloud architecture? How do we govern the eight different agent frameworks our business units have already adopted, without picking one as the official enterprise standard and crippling the other seven? What does FinOps look like for AI, when token costs and GPU-hours and training amortisation break the categories our finance organisation thought were settled? Above all: what is the architecture? Not the vendor list, not the platform standard — the architecture.

The frameworks the technology audience already uses are mature and structural. TOGAF organises their enterprise architecture practice. ITIL organises their service management. The FinOps Foundation framework organises their cloud financial governance. Platform engineering and SRE practices govern how they build and operate. Cloud well-architected frameworks — from AWS, Azure, and Google Cloud — organise their cloud deployment patterns. ISO/IEC 42001 is becoming the AI management system standard they will be audited against. These are operational frameworks; they describe how to manage architecture once it is drawn, but not the architecture itself for intelligent systems.

The OIA resonates with this audience because it is the architecture, drawn. It gives them a structural vocabulary for the conversations they are already trying to have with their boards and their CFOs. Layer 0 and Layer 1 are not abstractions when they are about GPU procurement and

silicon supplier concentration. Layer 2 is not abstract when it is about sovereign AI requirements and where regulated workloads can physically run. Layer 4 is not abstract when the question is whether to train internally, fine-tune a frontier model, or rely entirely on third-party inference. The OIA's distinction between layers that are typically partner territory (Layers 0, 1, 2, 4) and layers that are typically the enterprise's own stack (Layers 3, 6, 7, 8, 9) gives the CIO a framework for capital allocation conversations that has been missing.

Three parts of the OIA resonate strongly with the technology audience. First, Layer 8 — the Continuity Fabric — reframes 'agent state' from a framework-specific concern (which agent platform owns memory) to an architectural one (the enterprise must own continuity across whatever agent frameworks are in use, today and tomorrow). This relieves the CIO from the impossible position of having to pick a winning agent framework for the enterprise. Second, the cross-layer spans — particularly sovereignty (jurisdictional, operational, energy) and provenance — give the CIO an architectural way to articulate exposure that is currently being managed contractually, if at all. Third, the OIA's clarity about what is and is not a layer of the model lets the CIO push back on vendor positioning that conflates layers — a vendor claiming to provide 'the AI platform' is, by OIA reading, almost always claiming presence at four or five different layers, each of which deserves separate evaluation.

The CIO/CTO Lens organises the OIA against the operational frameworks the technology audience already uses. It is calibrated to the questions they are asking under capital pressure: build, buy, partner, license. It surfaces the layers where typical enterprises have de facto architecture but not de jure architecture, and gives the CIO a way to make that gap visible and governable. For the technology executive trying to articulate an AI architecture to a board that will hold them accountable for it, the CIO/CTO Lens is the document that makes the conversation tractable.

## **LENS 3 · THE GOVERNANCE AND RISK FUNCTION**

### **3.3 The Board and Risk Lens**

AI has moved from an operational concern to a board-level accountability faster than any technology category in recent memory. Directors who were comfortable overseeing cloud strategy a decade ago are now being asked to oversee something they understand structurally less well, with disclosure regimes evolving around them in real time, and with a fiduciary duty that requires more than nodding through the management presentation. The chief risk officer is being asked for an integrated view of AI risk that no existing risk taxonomy quite produces. The audit committee is being asked what to audit, in a category where the auditees are still figuring out what good looks like. The Board and Risk Lens is calibrated to those conversations.

The questions in board planning conversations in 2026 are about accountability and exposure. What is our material AI risk — and is the answer one we are prepared to disclose under SEC, FCA,

or equivalent regimes? Where does the EU AI Act apply to systems we are deploying or relying on, and have we classified them correctly? Who is accountable when an agent acts — the system owner, the model provider, the vendor, or the chief executive? What is our resilience posture if a model we depend on is silently changed by its provider, or if an agent in our environment behaves outside its intended scope? How is internal audit supposed to audit AI systems when the architectural vocabulary itself has not stabilised? And underneath all of these: are we, as a board, governing AI strategy — or are we receiving reports about it?

The frameworks already in their planning are increasingly specific. The NIST AI Risk Management Framework is becoming the lingua franca for AI risk governance and is being adopted internationally. ISO/IEC 42001 is the emerging AI management system standard the audit function will be asked to attest against. The EU AI Act is reshaping accountability for high-risk systems and general-purpose AI, with extraterritorial reach. SEC and equivalent disclosure regimes are pushing toward explicit articulation of material AI risks. Board fiduciary frameworks — Caremark in the United States, equivalent doctrines elsewhere — are being interpreted by the plaintiffs' bar in light of AI deployments. None of these frameworks is structural; they are accountability and disclosure frameworks. They tell the board what to oversee but not how the thing being overseen is built.

The OIA resonates with this audience because it provides the structural vocabulary the accountability frameworks assume but do not supply. When NIST AI RMF asks for evidence of governance across the AI lifecycle, the OIA's auditability and provenance spans tell the audit committee where in the architecture that evidence physically lives. When the EU AI Act requires that high-risk systems be subject to human oversight, the OIA's identity span and Layer 9 name the architectural surface at which oversight is implemented. When the SEC asks about material AI risks, the OIA's ten layers give the board a structured way to ask, layer by layer, where the exposure is concentrated and what the dependency map looks like — without forcing directors to become technologists.

Three properties of the OIA resonate especially strongly with directors and risk officers. First, the recognition that auditability cannot be retrofitted — it must be designed into every layer from the outset — is a clarifying observation for an audit committee that has been receiving 'we will improve auditability next quarter' as an answer for several quarters running. The OIA's framing makes that answer architecturally untenable. Second, the Continuity Fabric layer — named explicitly as the locus of recovery, rollback, and post-incident integrity for autonomous systems — maps directly to the resilience and disaster recovery questions board risk committees are accustomed to asking. Third, the OIA's three sovereignty dimensions — jurisdictional, operational, and energy — give the board a way to articulate sovereignty exposure as an integrated concern rather than as three disconnected ones.



The Board and Risk Lens treats the OIA's cross-layer spans — auditability, provenance, sovereignty, identity — as the primary axis of governance interest, with the layers themselves as the substrate against which exposure is mapped. This inversion (spans as primary, layers as substrate) is itself an example of the lens approach working: the same architecture, read differently for a different audience, without the architecture changing. The lens is built to support the conversations directors are actually having — with management, with regulators, with auditors, and with each other in executive session.

## 4. Future Lenses

---

The three published lenses cover the executive audiences most directly accountable for AI today. Other audiences sit just behind them, and the Lens System is designed to extend naturally to whichever ones the practitioner community needs next. The Foundation will support additional lenses where there is genuine demand and adequate practitioner engagement. The descriptions below outline the most likely candidates.

**The CFO Lens.** AI is the most capital-intensive enterprise technology investment of the past decade, and the categories of finance governance — CapEx versus OpEx, vendor concentration, working capital, unit economics — do not yet sit cleanly against AI deployment patterns. A CFO Lens would address capital efficiency under genuine uncertainty: how much AI spend is amortisable, how much is consumption-based, where supplier concentration creates contractual risk, and what the unit economics of an agent looks like at scale. Layers 0 through 2 of the OIA surface infrastructure and sovereignty as architectural concerns at the same level as application concerns — a CFO who can read the architecture by layer can govern AI capital allocation against a structural map rather than a vendor list.

**The General Counsel and Compliance Lens.** Liability, regulatory exposure, and contractual surface for AI deployment are evolving in real time. A General Counsel Lens would cover agent action liability, model output responsibility, training data IP exposure, vendor contract risk, EU AI Act conformity, US state-level AI law, data residency, and sectoral regulation in financial services and healthcare. The OIA's cross-layer spans — auditability, provenance, sovereignty, identity — are exactly the artefacts that legal accountability and regulatory compliance reduce to. A general counsel can read the OIA and locate, layer by layer, where contractual and regulatory exposure physically lives.

**The CHRO Lens.** Workforce transformation is the half of the AI conversation that has been least well served by existing frameworks, because the existing frameworks were built for human workforces only. A CHRO Lens would address AI augmentation versus displacement, the governance of agents as a non-human workforce alongside human employees, performance

management in mixed teams, reskilling at scale, and the increasingly serious question of how to onboard, scope, behaviorally baseline, and offboard agents in the same governance discipline applied to people. Layer 7, Layer 9, and the identity span of the OIA together describe the architectural surface at which workforce policy for agents is implemented.

Other lenses are possible — a chief data officer lens, a regulated-industry lens for financial services or healthcare, a public-sector lens, an operating-executive lens for a line of business with significant AI exposure. The Foundation will publish lenses for any audience where the work meets the discipline of the approach. Practitioners and organisations with deep audience expertise are invited to contribute proposals for new lenses; the Foundation will collaborate on authoring, review, and publication.

These additional lenses are described here as future possibilities, not commitments. The current release publishes the three lenses for the audiences with the deepest need and the most developed framework vocabulary. Subsequent lenses will follow as practitioner engagement and demand support them.

## Closing

---

The OIA Reference Model describes the architecture of enterprise intelligent systems with the rigour the discipline requires. The OIA Lens System is how that architecture reaches the executives who must apply it — each in their own vocabulary, against their own frameworks, around their own questions. Together, the architecture and the lenses make the model usable across the executive team that will collectively own AI as an enterprise concern.

The Foundation invites the executive audiences described in this document to engage with the OIA Reference Model directly, with the published CISO Lens, CIO/CTO Lens, and Board and Risk Lens as the entry points calibrated to their respective disciplines. The Foundation also invites practitioners with deep expertise in audiences not yet represented — finance, legal and compliance, human resources, and others — to contribute to the lens roster as it grows. Reference architectures earn their place when the people who must apply them can find their own work in the language. The Lens System is the path to that recognition.